

Component Interaction Matrix				
Component	Input Data	Processing	Output Data	Downstream
Layer 2: Validation	Raw fifty-six features, 15,120 obs (elevation, slope, aspect, etc.)	Validation: range, types, completeness;	Valid fifty-six features, 15,120 obs (nulls flagged/imputed, metadata timestamps, validation report)	Feature Engineering (Layer 3)
Layer 3A: Elevation Processing	Elevation (continuous) 1859–3858 m raw elevation per obs	<ul style="list-style-type: none"> <li>- Bin into elevation zones (<math>\leq 2400</math>, 2400–2800, 2800–3200, <math>&gt; 3200</math>)</li> <li>- Detect proximity to thresholds (<math>\pm 50</math> m windows)</li> <li>- Produce distance-to-threshold metric</li> </ul>	Elevation zone categorical  Threshold flags (near_2400, near_2800...)  distance_to_nearest_threshold	Model Training  Feature Combin.
Layer 3B: Aspect Transform	Aspect (0–360°), Slope (cont.) (slope 0–??°)	<ul style="list-style-type: none"> <li>- Convert aspect <math>\rightarrow</math> <math>\sin(\text{aspect})</math>, <math>\cos(\text{aspect})</math></li> <li>- Normalize slope (robust scaler / quantile)</li> </ul>	$\sin(\text{aspect})$ , $\cos(\text{aspect})$ normalized_slope	Feature Combin.  Model Training

		- Outlier smoothing for slope		
Layer 3C: Soil Consolidation	Forty soil types (sparse categorical) one-hot / indexes per obs	- Consolidate soil types by frequency + ecological similarity - Create fallback bucket "rare_soil." - Version mapping rules (config JSON)	Fifteen consolidated soil feature columns (one-hot) soil_group_id, soil_group_desc	Feature Combin.  Model Training
Layer 3D: Distance Interact	Distances (4 types), Hillshade (3) (to hydrology/road/fire points), hillshade_9am/noon/3pm	- Normalize distances (log / min-max) - Combine distances → interaction features (e.g., ratio, product) - Generate hillshade deltas and temporal features	normalized_distances (4) interaction_features (≈6) hillshade_features	Layer 4 Training
Layer 4: Model Training	35–40 Engineered Features (train/val) labels (7 cover classes) hardware spec (CPU/GPU)	- Train base models separately: RF, XGBoost, LightGBM - Hyperparam search (Optuna, one hundred trials)	3 Base models (artifacts) per-model CV metrics, feature importances training logs, artifacts in MLflow	Ensemble Integrat. (Layer 4.2)

		<ul style="list-style-type: none"> <li>- CV strategy: stratified (BUT MUST use spatial blocked CV)</li> <li>- Calibration step (Platt/Isotonic) opcional</li> </ul>	inference time per model	
Layer 4.2: Ensemble Integr.	3 Base model artifacts + CV preds model metrics, latency constraints	<ul style="list-style-type: none"> <li>- Ensemble: weighted voting (w= [0.3,0.4,0.3])</li> <li>- Option: stacking (meta-learner) for experimental runs</li> <li>- Export distilled model (optional)</li> </ul>	<ul style="list-style-type: none"> <li>- Ensemble artifact (v1.0), ensemble weights</li> <li>- Ensemble CV performance (accuracy 95.2%)</li> <li>- export formats: ONNX / native boosters</li> </ul>	Layer 5 Prediction & Uncertainty
Layer 5: Prediction & Unc.	Ensemble outputs (per-model probs) per obs; feature metadata; thresholds	<ul style="list-style-type: none"> <li>- Combine probabilities: weighted average → final prob vector</li> <li>- Compute aleatoric uncertainty (entropy <math>H(p)</math> normalized by <math>\log_2</math>)</li> <li>- Compute epistemic uncertainty (variance across models/ensem)</li> </ul>	Prediction JSON per obs: <ul style="list-style-type: none"> <li>- cover_type (argmax)</li> <li>- probability distribution (7)</li> <li>- aleatoric, epistemic, total_uncertainty</li> </ul>	Layer 6 Monitoring & Alerts / Decisioning

		<ul style="list-style-type: none"> <li>- total_uncertainty = sqrt (aleatoric^2 + epistemic^2)</li> <li>- Apply threshold proximity rules: amplify uncertainty ×2 (configurable)</li> <li>- Generate warnings: e.g., "threshold_proximity"</li> </ul>	<ul style="list-style-type: none"> <li>- confidence_score = 1 - total_uncertainty (capped)</li> <li>- warnings list (text, recs)</li> </ul>	
Layer 6: Monitoring & Drift & Sensitivity	Stream of predictions + input stats (elevation hist, feature distributions) baseline metrics (training distrib.)	<ul style="list-style-type: none"> <li>- Drift detection (KL-divergence per feature, PSI, KS tests)</li> <li>- Threshold window counters (±50 m counts &amp; % near thresholds)</li> <li>- Alerting rules (minor/major drift, %accuracy drop)</li> <li>- Trend analysis (accuracy week1→week4 degradation)</li> <li>- Risk scoring for manual review prioritization</li> </ul>	<ul style="list-style-type: none"> <li>Drift reports (KL per feature, PSI)</li> <li>Threshold windows: counts, % obs near windows.</li> <li>Alerts (Grafana / PagerDuty)</li> <li>Alerts (Grafana / PagerDuty)</li> <li>Drift_action_suggestion (retrain/canary/manual)</li> </ul>	Ops / Retraining (canary, retrain) Field verification

<p>Layer 7A:</p> <p>REST API</p> <p>Serving (real-time)</p>	<p>Single-observation request (JSON) (56 Features)</p>	<ul style="list-style-type: none"> <li>- API ingress (NGINX) → FastAPI/Uvicorn</li> <li>- Auth, schema validation, routing</li> <li>- Retrieve model artifact from model cache (Redis) or local mem</li> <li>- Run inference, post-process (warnings, uncertainty)</li> <li>- Log to Postgres (predictions_log)</li> <li>- Return response</li> </ul>	<p>API Response</p> <p>JSON:</p> <ul style="list-style-type: none"> <li>- prediction (cover_type, name)</li> <li>- probabilities (7), confidence_score</li> <li>- uncertainty breakdown (aleatoric/epistemic)</li> <li>- metadata: model_version, latency_ms</li> <li>- warnings &amp; recs p50/p95/p99 latency metrics</li> </ul>	<p>Client apps / GIS</p> <p>Field apps, dashboards</p>
<p>Layer 7B:</p> <p>Batch Processing</p>	<p>Batch CSV / GeoTIFF (N samples) (10k → 1M samples)</p>	<ul style="list-style-type: none"> <li>- Batch orchestration (K8s jobs / Spark / GPU-batching)</li> <li>- Parallel GPU-accelerated inference; chunking &amp; checkpointing</li> <li>- Store results to S3 (batch results + uncertainty maps)</li> </ul>	<p>Batch outputs:</p> <p>CSV/GeoTIFF/JSON</p> <ul style="list-style-type: none"> <li>- prediction per row, uncertainty maps (GeoTIFF)</li> <li>- processing logs &amp; performance report</li> </ul>	<p>Policy teams / GIS</p> <p>Reforestation planning</p>

		- Emit metrics to Grafana & update DB for audits	- summary stats (coverage by class, %near thresholds)	
Model Registry (MLflow)	Model artifacts (trained), metadata. (hyperparams, training logs)	<ul style="list-style-type: none"> <li>- Register model versions (MLflow): metrics, artifacts, tags</li> <li>- Promote (staging → production) after tests</li> <li>- Store FE_version and data lineage links</li> </ul>	MLflow entries: model_id, version, stage  Link to artifacts (S3/registry)	Serving / Auditability CI/CD pipelines
Model Cache (Redis)	Serialized serving artifacts (hot) (ONNX, native boosters)	<ul style="list-style-type: none"> <li>- Keep warmed model for low-latency loads (LRU eviction)</li> <li>- Heartbeat checks; reload on version change</li> </ul>	Fast model load (mem snapshot)  cache_hit/miss metrics	API Serving (Layer7A)
Storage (S3)	Trained models, batch results, maps GeoTIFFs, uncertainty maps	<ul style="list-style-type: none"> <li>- Versioned object storage; lifecycle rules (archive)</li> <li>- Serve as MLflow artifact store</li> </ul>	Artifacts accessible for batch & registry  Long-term archive, access logs	Audits / Reproducibility
Database (Postgres)	Predictions logs, metrics, alerts	<ul style="list-style-type: none"> <li>- Store transaction logs, user queries, monitoring snapshots</li> </ul>	Queryable logs: predictions_table, metrics	Reporting / Audits

		- Index for fast retrieval; retention & partitioning policies	alerts_table (drift, threshold hits)	
Monitoring (Grafana) & Alerting	Metrics (latency, accuracy, drift) traces (Tempo), logs	- Dashboards & alerting rules; integrate with Prometheus - Alert escalation (PagerDuty), runbooks link	Visual dashboards: latency p50/p95/p99, drift Alert tickets, on-call pages	SRE / ML Ops